

# AlphaEarth Foundations

Explainer & Implementation by Brayden Zhang

# Summary


- AEF is a geospatial foundation model that produces embedding fields, which can then be used for downstream geographical/environmental applications
- Key Features
  - Multisource integration
  - Space-Time-Precision Encoder
  - Time-conditional summarization
  - vMF bottleneck for embeddings
  - Multi-term loss: reconstruction, batch uniformity, teacher-student consistency, CLIP alignment.

# Data

- Data Types
  - Sentinel-2 (optical, 13 bands, 10–60m).
  - Sentinel-1 (radar, dual-pol).
  - Landsat 8/9.
  - ERA5 climate reanalysis.
  - GEDI LiDAR (vertical canopy profiles).
  - Ancillary labels: Wikipedia text.
- Each training sample:
  - Support period: multi-frame inputs (up to 1 year).
  - Valid period: target interval for summarization.
  - Spatial patch: e.g., 128x128 px.
  - Output: per-pixel embeddings (B,H,W,64).

# Encoder

D




# Space Operator



- **Goal:** to capture long-range spatial dependencies inside each frame
- Input:  $x \in \mathbb{R}^{B \times T \times H \times W \times C}$
- Flatten Spatial Grid  $x_{\text{flat}} \in \mathbb{R}^{(BT) \times (HW) \times C}$   
 $[Q, K, V] = x_{\text{flat}} W_{qkv}, \quad W_{qkv} \in \mathbb{R}^{C \times (3 \cdot C)} \quad Q, K, V \in \mathbb{R}^{(BHW) \times \text{heads} \times T \times d}$
- Perform (per-frame) MHSA
- Residual + MLP  $\text{Attn}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$
- Output: (B, T, H, W, C) at 1/16L resolution

# Time Operator



- Goal: to capture temporal evolution of each pixel over many frames
- Input:  $x \in \mathbb{R}^{B \times T \times H \times W \times C}$
- Add Sinusoidal temporal encodings  $x_{b,t,h,w} \leftarrow x_{b,t,h,w} + \text{TimeEnc}(t)$
- Flatten along time at each pixel location
$$x_{\text{flat}} \in \mathbb{R}^{(BHW) \times T \times C} \quad [Q, K, V] = x_{\text{flat}} W_{qkv} \quad Q, K, V \in \mathbb{R}^{(BHW) \times \text{heads} \times T \times d}$$
- Temporal MHSA
- Residual + MLP 
$$z = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$
- Output: (B, T, H, W, C) at 1/8L resolution

# Precision Operator



- Goal: to be able to preserve finer local detail with Convolutions
- Input  $x \in \mathbb{R}^{B \times T \times H \times W \times C} \implies (BT, C, H, W)$
- Apply 3x3 convolutions

$$x' = \text{Conv}_{3 \times 3}(\text{GroupNorm}(x))$$

- Expansion + Contraction Conv + Residual

$$x_{\text{exp}} = \text{GELU}(\text{Conv}_{3 \times 3}^{C \rightarrow 4C}(x_{\text{norm}})) \quad x_{\text{contr}} = \text{Conv}_{3 \times 3}^{4C \rightarrow C}(x_{\text{exp}}) \quad x' = x + x_{\text{contr}}$$

- Output at L/2 resolution

$$x' \in \mathbb{R}^{B \times T \times H \times W \times C}$$

# Pyramid Exchanges



- Purpose: a general way to *share learnings* across S/T/P operators
- At the end of each STP block, each operator has features at its native resolution
  - Precision (B,T,L/2,L/2,Cp)

The *learned Laplacian pyramid resampler* takes each output and resizes it to the resolutions of the other operators

- Precision → downsampled to L/8, L/16
- The resampled outputs are fused and sent into the next block

# Implementation

```
class LearnedSpatialResampling(nn.Module):
    """Learned Laplacian pyramid rescaling for spatial pyramid exchanges."""

    def __init__(self, in_channels: int, out_channels: int, scale_factor: float):
        super().__init__()
        self.scale_factor = scale_factor


        if scale_factor > 1:
            # Upsampling: scale_factor > 1 means output is larger
            # For scale_factor = 2.0, output size = input_size * 2
            self.conv = nn.ConvTranspose2d(
                in_channels, out_channels,
                kernel_size=4, stride=2, padding=1
            )
        elif scale_factor < 1:
            # Downsampling: scale_factor < 1 means output is smaller
            # For scale_factor = 0.5, output size = input_size / 2
            stride = int(1 / scale_factor)
            self.conv = nn.Conv2d(
                in_channels, out_channels,
                kernel_size=stride * 2 - 1,
                stride=stride,
                padding=stride - 1
            )
        else:
            # Same resolution
            self.conv = nn.Conv2d(in_channels, out_channels, kernel_size=1)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.conv(x)
```

self.space\_to\_time = LearnedSpatialResampling(self.space\_dim, self.time\_dim, 2.0)  
self.space\_to\_precision = LearnedSpatialResampling(self.space\_dim, self.precision\_dim, 8.0)  
self.time\_to\_space = LearnedSpatialResampling(self.time\_dim, self.space\_dim, 0.5)  
self.time\_to\_precision = LearnedSpatialResampling(self.time\_dim, self.precision\_dim, 4.0)  
self.precision\_to\_space = LearnedSpatialResampling(self.precision\_dim, self.space\_dim, 0.125)  
self.precision\_to\_time = LearnedSpatialResampling(self.precision\_dim, self.time\_dim, 0.25)


# Summary Period Encoder

- After the STP encoder, we still have a time-series of features per pixel
- Objective: to reduce the sequence of features to a single embedding per pixel that summarizes a chosen valid period  $[t_s, t_e]$
- Valid Period Encoding: encode start, end, and duration with sinusoidal + MLP features
- Query Vector: learned query conditioned on valid-period encoding
- Attention pooling: have the query attend over temporal features at each pixel location
- Outputs: Embedding Grid ( $B, H, W, 64$ )




# Von Mises-Fisher Bottleneck

- Constrain embeddings to the unit hypersphere
- Each embedding is modeled as the mean direction of a von Mises-Fisher distribution (a probability distribution on the hypersphere  $S^{d-1}$ , *in our case d = 64*)
- Project STP features → 64D vector
- Apply L2norm
- Output: per-pixel unit vector  $\mu \in S^{63}$



# Source Decoders

- Purpose: to reconstruct raw observations from embeddings; used for training
- Lightweight decoders per source
- Using MLPs
- Input ( $B, H, W, 64$ )
- Outputs:
  - Continuous values or Categorical classes



# Loss Function


Loss = Reconstruction Loss + Batch Uniformity Loss + Teacher Student Consistency Loss +  
Text Contrastive (CLIP) Loss

$$l = \frac{a}{M} \sum_{i \in M} f_i(\mathbf{y}_i, \mathbf{y}'_i) w_i + b \sum_{i=1}^{64} |u_i \cdot u'_i| + c \left( \frac{1 - \mathbf{u} \cdot \mathbf{u}_s}{2} \right) + d f_{\text{CLIP}}(\mathbf{u}, \mathbf{u}_t) \quad (3)$$

# Reconstruction Loss

- Compares the predicted observation  $y_i'$  to the true observation  $y_i$  for each source  $i$
- Objective is to force the embeddings to carry enough info to be able to reconstruct the raw inputs
- Using L1 loss
- Weights  $w_i$  allow scaling importance by source


$$a \cdot \sum_i f_i(y_i, y'_i) \cdot w_i$$



RECONSTRUCTION

# Batch Uniformity Loss

- Objective: to make the embeddings (64D) to be uniformly distributed
- Take an embedding batch  $\mu_1, \mu_2, \dots, \mu_B$
- Rotate it by 1 along the batch dimension:
- $\mu_{rot} = [\mu_B, \mu_1, \mu_2, \dots, \mu_{B-1}]$
- Pair each original vector with the rotated one:  
 $(\mu_1, \mu_B), (\mu_2, \mu_1), \dots$



$$b \cdot \text{BatchUniformity}(\mu)$$

Compare the dot product for each pair:


$$L_{uniform} = \frac{1}{B} \sum_{i=1}^B |\mu_i \cdot \mu_{rot,i}|$$

- If the embeddings are aligned (dot product closer to 1)  $\rightarrow$  loss is larger
- If the embeddings are further apart (dot product closer to 0)  $\rightarrow$  loss is smaller

```
x = torch.nn.functional.normalize(x, p=2, dim=-1)
# Rotate (roll) sample pairs to approximate u' in the paper
x_prime = torch.roll(x, shifts=1, dims=0)
dots = (x * x_prime).sum(dim=-1).abs() # |u · u'|
```

# Teacher-Student Consistency Loss

$$c \cdot \frac{1 - \mu \cdot \mu_s}{2}$$



$\mu$  = teacher embedding (all data sources)

$\mu_s$  = student embedding (only a few data sources)

Dot product measures cosine similarity so if the student matches the teacher:

dot product = 1

loss = 0

# Text Contrastive Loss

$$d \cdot f_{\text{CLIP}}(\mu, u_t)$$

```
def clip_loss(self, image_embeddings: torch.Tensor,
              text_embeddings: torch.Tensor) -> torch.Tensor:
    """Compute CLIP-style contrastive loss."""
    # Expect (B, D) vs (B, D)
    img = torch.nn.functional.normalize(image_embeddings, p=2, dim=-1)
    txt = torch.nn.functional.normalize(text_embeddings, p=2, dim=-1)
    logits = img @ txt.t()  # (B, B)
    targets = torch.arange(img.size(0), device=img.device)
    loss_i = torch.nn.functional.cross_entropy(logits, targets)
    loss_t = torch.nn.functional.cross_entropy(logits.t(), targets)
    return 0.5 * [loss_i + loss_t]
```